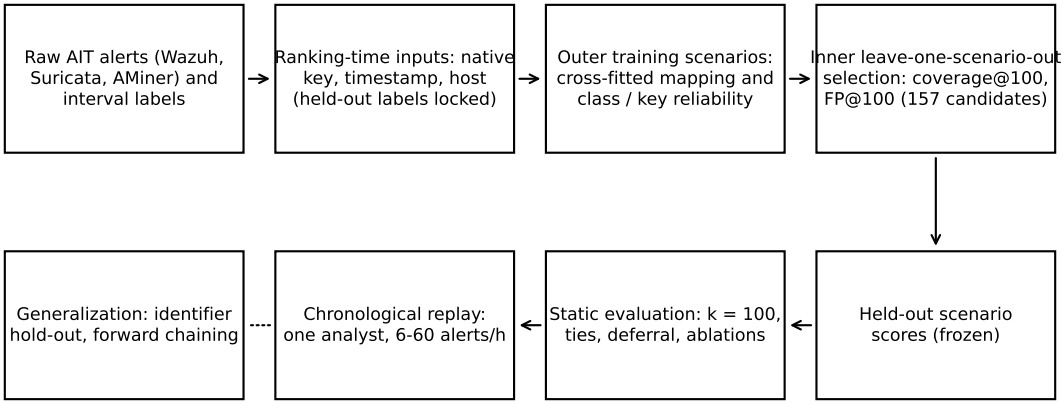
Raw AIT alerts (Wazuh, Suricata, AMiner) and interval labels
Ranking-time inputs: native key, timestamp, host (held-out labels locked)
Outer training scenarios: cross-fitted mapping and class / key reliability
Inner leave-one-scenario-out selection: coverage@100, FP@100 (157 candidates)
Held-out scenario scores (frozen)
Static evaluation: k = 100, ties, deferral, ablations
Chronological replay: one analyst, 6-60 alerts/h
Generalization: identifier hold-out, forward chaining